

Practical Aggregate Programming with Protelis

Danilo Pianini ¹ Mirko Viroli ¹ Jacob Beal ²
{danilo.pianini, mirko.viroli}@unibo.it
jakebeal@ieee.org

¹ALMA MATER STUDIORUM—Università di Bologna

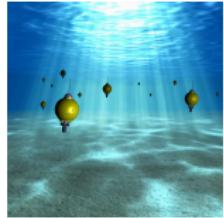
²Raytheon BBN Technologies

SASO Tutorial
2017-09-22 – Tucson, AZ, USA

The IoT is becoming a crowded and complex place

Future and emerging Internet-of-things are witnessing:

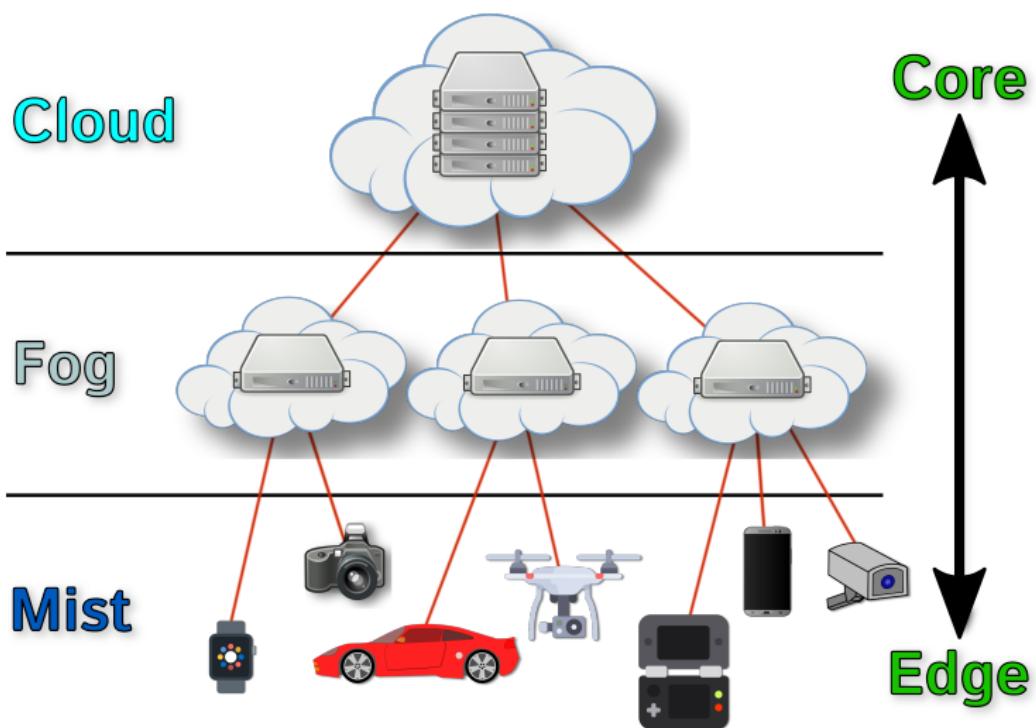
- devices increasingly wearable, mobile, embedded, flying...
 - increasing availability of heterogeneous wireless connectivity
 - increasing availability of computational resources (device/edge/cloud)
 - increasing production and analysis of data, everywhere, anytime
- ⇒ business, security, and privacy concerns will probably be drivers as well



Programming models for mobile and IoT applications

- Client side
 - Single-device program: objects, functions, actors, futures, tasks, activities
 - Local intelligence
- Client-server communication
 - SOA (Http/Rest), MoM (RabbitMQ)
- Server side
 - Storage: OO, relational, NoSQL — in memory or on-disk
 - Coordination (orchestration, mediation, rules enactment)
 - Situation recognition (online/offline, mining, business intelligence, stream processing)
- Scalability in the server calls for cloud-ification
 - Not really orthogonal to the whole programming model
 - It often dramatically affects system design

Heterogeneity of deployment contexts



Implications

Where does programming effort end up into?

Programs of clients and servers highly depend on:

- Platform and API
- Communication technology
- Number, type, and dislocation of devices

⇒ Design and deployments hardly tolerate changes

⇒ IoT systems tend to be rigid, and costly to debug, evolve, and maintain

The technological result

- Few of the opportunities of large-scale IoT are currently exploited
 - Virtually, any computational mechanism (sensing, actuation, processing, storage) could involve spontaneous, adaptive cooperation of large sets of devices
- How many large-scale deployments of **adaptive** IoT systems around?
- Where are the “Collective Adaptive Systems”?

1 Introduction

- Aggregate programming: idea and motivation
- Protelis as practical AP language and workbench setup

2 Core language constructs

- The boring stuff: functions, operators, and the like
- Stateful computation with `rep`
- Sharing information with `nbr`
- Domain restriction with `if`

3 Libraries

- Frequently used patterns: towards reusable building blocks
- Filling the abstraction gap: Protelis-lang

4 More about the toolchain

- Making nodes move
- Configuring batches and exporting data
- Using real world maps
- Getting from simulations to devices

Outline

1 Introduction

- Aggregate programming: idea and motivation
- Protelis as practical AP language and workbench setup

2 Core language constructs

- The boring stuff: functions, operators, and the like
- Stateful computation with `rep`
- Sharing information with `nbr`
- Domain restriction with `if`

3 Libraries

- Frequently used patterns: towards reusable building blocks
- Filling the abstraction gap: Protelis-lang

4 More about the toolchain

- Making nodes move
- Configuring batches and exporting data
- Using real world maps
- Getting from simulations to devices

A programming model perspective

What do we lack in large-scale IoT systems?

- The plain old platform-independent programming abstraction
- Delegating to the underlying platform everything but the core logic:
 - non-functional aspects (network performance, battery saving...)
 - deployment issues

The challenge

Let's tackle the worst possible scenario:

- Countless heterogeneous, moving devices unpredictably situated
- Pervasive sensing and actuation
- Contextual (spatio-temporal) behavior
- Abstract away the infrastructure
 - We can't rely on the cloud or fog
 - But we must be ready to exploit the opportunity they provide

Towards aggregate computing I

Desiderata

- Abstractions well suited to the scenario
- Techniques that promote solid engineering principles
 - Reusability
 - Composability
 - Modularity
- A toolchain for designing, programming, and testing
- As few constructs as possible, tractable semantics

Towards aggregate computing II

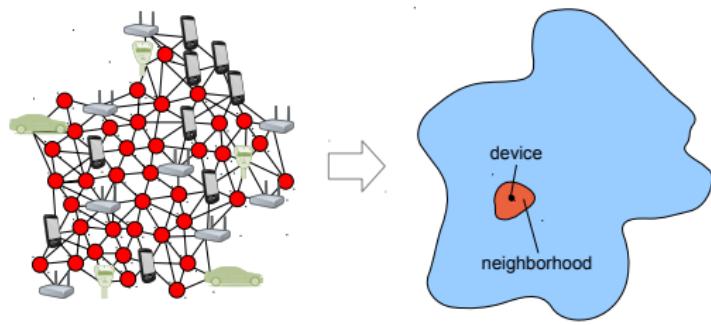
Aggregate computing

- Based on the “field calculus”: a core, tractable calculus
- Relies on the idea of “computational field”
- Similar to functional languages
- Promotes programming as “describing the problem” rather than “hacking a solution”
- Practical languages based on such calculus are available...
- ...as well as simulation environments where to perform testing
- Somehow very different from languages rooted on other paradigms

Manifesto of aggregate computing

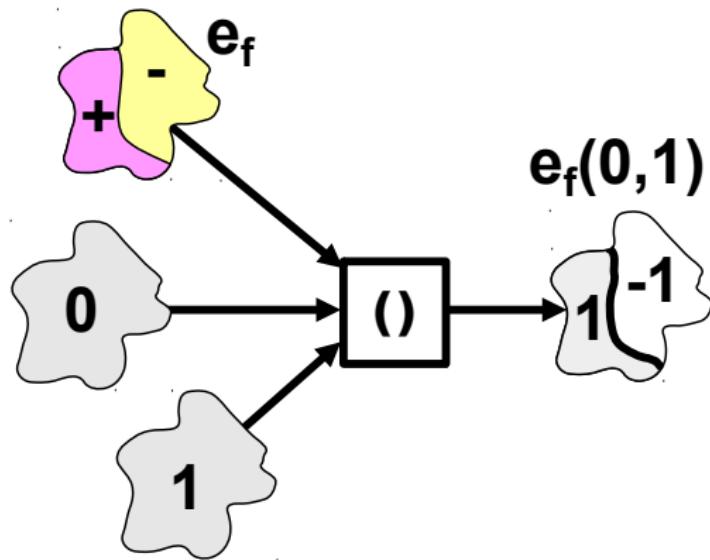
Motto: program the aggregate, not individual devices!

- ① The reference computing machine
⇒ an aggregate of devices as a single entity fading to the actual *space*
- ② The reference elaboration process
⇒ atomic manipulation of a collective data structure
- ③ The actual networked computation
⇒ a proximity-based self-org system hidden “under-the-hood”



Computing with fields

Field: a distributed data structure mapping each device to a value



A look under the hood

The computational round semantics

Every device operates as follows:

- ① Sleeps from some time
- ② Wakes up, discards the old information, processes the new information it may have received, and computes its program
- ③ Sends the results away, and goes back to sleep

Synchronization

- No synchronization required
- Devices should compute rounds at a similar frequency
- The impact of very diverse frequencies depends on the specific program

Protelis [PVB15]: aggregate programming in practice

- The Field Calculus is a “theoretical object”
 - It can (and is) used for mathematical proofs, but provides no compiler or runtime
- Protelis is an attempt at making the higher order field calculus practically usable through:
 - An executable interpreter
 - A concrete C-family syntax
 - Interoperability with the Java Runtime
 - With access to the existing plethora of Java libraries
 - A set of tools for development, including Eclipse plugins

Running Protelis programs

- Protelis is to be used to program aggregate of devices
- We need to run the programs on a network of devices
 - A single device is a very special case of a network of devices...
 - ...but a purely local run does not normally provide any useful information about the algorithm
- We need to run the programs on a network of devices
 - Needless to say how impractical is to deploy a network every time you want to perform a run
- Another tool is required!
 - Protelis exposes an API for writing back-ends in charge of device, hardware, and network interaction
 - It's rather easy to make Protelis work upon existing frameworks...
 - ...hereby including simulated environments

Alchemist as execution platform for Protelis

Alchemist [PMV13] is a simulator inspired by stochastic chemistry

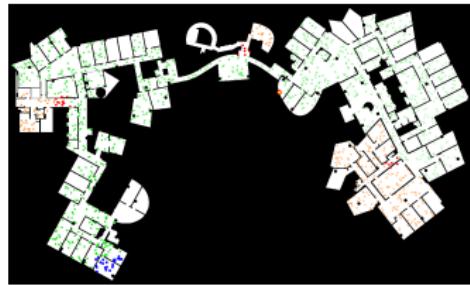
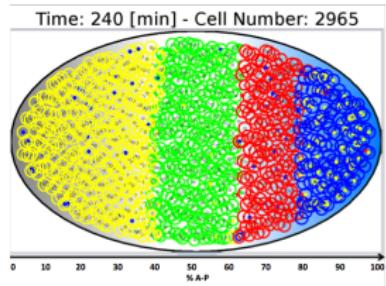
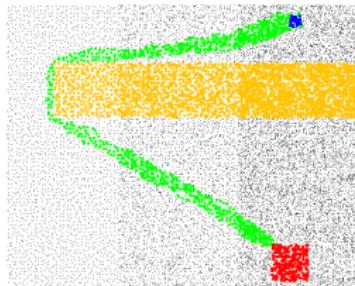
Features (the subset of interest for this tutorial)

- Support for executing Protelis programs
- Several types of environments:
 - Euclidean spaces, possibly with obstacles (for indoor environments)
 - Real world maps (using Openstreetmap data), support for GPS traces
- YAML [BKEdN01] as specification language
- Multiplatform (hosted on Java 8+)

Already in use for

- Crowd evacuation [MVR⁺11]
- Crowd steering [BPV15]
- Pervasive computing [MPV12a, SYD⁺13, PBV16, BVPD16, VBPB16, PDV17]
- Biochemistry [MPV12b]

Simulation examples



Workbench setup I

Let's have fun together!

Experimenting is the best way of learning a new paradigm and/or language

Prerequisites

- Internet connection
 - Though I might have prepared a stand-alone version on a USB drivers
- A working installation of Java 8+
- A terminal emulator
 - Unix-like terminals are ok
 - Should work also with the Windows cmd.exe
 - Don't ignore case sensitivity
- git (preferred), or a program to decompress a zip file

Workbench setup II

Download

- In the following, WORKSPACE represents the folder you want to work into
- **Option 1 – git is installed on the system. Issue:**
 - git clone <https://github.com/AlchemistSimulator/SAS017Tutorial.git> WORKSPACE
 - cd WORKSPACE
- **Option 2 – git???**
 - Download from: <http://bit.ly/saso17tutorial-zip>
 - Unzip in WORKSPACE
 - Open a terminal in WORKSPACE

Workbench setup III

First run

Issue the following command:

- On Unix terminals:
 - `./gradlew -Pfile=nothing`
- On Windows Prompt:
 - `gradlew.bat -Pfile=nothing`

Expectation:

- Lots of downloads
- More downloads
- A graphical interface with a white background

Setup explained I

The Gradle mini-box

Gradle is a (very trendy) build system based on Java and Groovy. We use it to:

- Download the software automagically
- Execute the simulator
 - It's easier, quicker and more portable than using sh or java directly
- It reads what to do from the `build.gradle` file

Setup explained II

build.gradle

```
// We will need Java
apply plugin: 'java'

// All the stuff can be found on Maven Central
repositories { mavenCentral() }

// We need the simulator at version 6.0.2 (change the number to test with newer versions!)
dependencies { compile "it.unibo.alchemist:alchemist:7.0.0" }

// OK that's a bit more complicated, but in the end just configures a Java process and launches it
task "run$file"(type: JavaExec){ // Create a new task with a dynamically discovered name (Groovy coolness)
    classpath = sourceSets.main.runtimeClasspath // Put the simulator and its dependencies in the classpath
    classpath 'protelis' // Put all the protelis files in the classpath (so the interpreter can find them)
    classpath 'maps' // Cool stuff for the last example :)
    main = 'it.unibo.alchemist.Alchemist' // Launch this java class
    args(
        "-y", "sim/${file}.yml",
        "-g", "effects/${file}.aes",
        "-e", "exported-data"
    ) // With parameters "-y sim/<SOMETHING>.yml -g effects/<SOMETHING>.aes -e exported-data"
}

/*
 * If nothing different is specified, search for a property named "file" in the project, then run a task
 * with such name.
 */
defaultTasks "run$file"
```

A simulated environment

Let's displace 500 nodes in a circle in a bidimensional euclidean space.

00-devices-in-a-circle.yml

```
incarnation: protelis
displacements: # In this section we list our device displacements
  - in:
      type: Circle
      parameters:
        - 500 # Number of devices
        - 0 # X center of the circle
        - 0 # Y center of the circle
        - 50 # radius
```

- Run with `./gradlew -Pfile=00-devices-in-a-circle`
- Expected: 500 devices (black dots) randomly displaced inside a circle

Pan, zoom, node information

Let's play a little bit with the graphical interface:

- Pan by left click + drag
- Zoom with the mouse wheel
- Rotate by right click + drag

The node closest to the mouse is marked with a yellow and red dot

- Double click to open a frame with the node information
 - Position
 - Content
 - Programs to execute
- Press **M** on your keyboard to disable the marker
 - Useful e.g. if you want to grab a screenshot

Manually moving nodes around

Nodes can be moved around by hand:

- The **S** key enters and exits the “selection mode”
- Enter in selection mode, select a group of nodes by dragging a rectangle around them
- Press **O**
- Drag the nodes to another position

Be wary that:

- Moving around is very useful for demoing and prototyping
- Moving stuff manually around will make your experiment **non reproducible**
- We'll see other means of moving nodes and retain reproducibility

Drawing on screen in Alchemist I

Alchemist supports the definition of a stack of graphic effects in order to use colors and shapes to understand what's going on.

- Saved as a JSON file, that can be manually edited before launching the simulation
- The effects can be edited and saved from within the GUI
- The stack is applied one effect at a time
 - Upper layers cover lower layers
 - With some training, it's possible to achieve nice results using transparencies
- The effect system, and the whole GUI, is on a process of heavy reworking
 - We are switching to the modern JavaFX
 - We are making the effects much easier to create and manipulate from the UI
 - The plan is to release it by the end of 2017, stay tuned

Drawing on screen in Alchemist II

00-devices-in-a-circle.aes

```
[  
 {  
   "type": "class it.unibo.alchemist.boundary.gui.effects.DrawShape",  
   "curIncarnation": "protelis",  
   "mode": "FillEllipse",  
   "red": {  
     "max": 255,  
     "min": 0,  
     "val": 0  
   },  
   "blue": {  
     "max": 255,  
     "min": 0,  
     "val": 0  
   },  
   "green": {  
     "max": 255,  
     "min": 0,  
     "val": 0  
   },  
   "alpha": {  
     "max": 255,  
     "min": 0,  
     "val": 255  
   },  
   "scaleFactor": {  
     "max": 100,  
     "min": 0,  
     "val": 50  
   }  
 },  
 ]
```

Network model

If we want the node to connect to each other, we must tell the simulator how to do so:

01-devices-connected.yml

```
incarnation: protelis
# Connect nodes that are sufficiently close
network-model:
  type: EuclideanDistance
  parameters: [10]
displacements:
  # Expected: a YAML list of device displacements
  - in:
    type: Circle
    parameters: [500, 0, 0, 50]
```

- Launch with `./gradlew -Pfile=00-devices-connected`
- Expected: same as the previous experiment
- Press **L** on the keyboard to visualize the network links!
- Try to move them (**S** then **O** then drag and drop)

On the type ... parameters syntax

Alchemist can load any implementation of its interfaces

For each type of object written in the simulation file, Alchemist tries to find and build a Java object implementing a concrete version. The resolution is operated as follows:

- ① if type and (optionally) parameters are found as keys of the object, a class with the corresponding name is searched and a constructor is invoked with the provided parameters
 - All parameters that can be deduced by the context are automatically bound (objects of type Environment, RandomGenerator, Incarnation, Node, Reaction, LinkingRule) and must not be provided in the parameters list
- ② If the provided parameter is a string, it is forwarded to the current Incarnation for interpretation
- ③ In any other case, a default system value is looked for

If the lookup or the build fail, the simulation won't get built.

To find out the available types and their constructors, refer to the Javadoc at:
<http://alchemist-doc.surge.sh/>

My first Protelis program I

- Now press **P** to begin the simulation

Uh?! Nothing happens... The time just jumps to Infinity.

- In fact, there are no events in this simulation :)

Let's program the nodes with an extremely simple Protelis program: 0

My first Protelis program II

02-0.yml

```
incarnation: protelis
network-model:
  type: EuclideanDistance
  parameters: [10]
displacements:
- in:
  type: Circle
  parameters: [500, 0, 0, 50]
programs:
- # This list of lists seems useless, it actually makes sense when you use YAML anchors
  # Round frequency (advanced distributions can be used with the time/parameters syntax)
- time-distribution: 1
  # This is the program we want to execute locally
  program: 0
  # This tells the simulator to share the results of the last computation. No time distribution has
  ↪ ``as soon as possible'' semantics (Markovian event with infinite rate). The separation is
  ↪ useful in case the user wants to simulate a delay between the computation and the delivery
  ↪ of messages, e.g. to simulate network delays, or a battery saving policy on the wireless
  ↪ modules.
- program: send
```

- Run with: `./gradlew -Pfile=02-0`
- Then press **P** to begin the simulation

My first Protelis program III

Expected:

- The time flows!
- Nodes turn red
 - I've configured an effect that changes color if the node has computed, it's no black magic
- If you open a window with the node contents:
 - In "Contents", 0 is associated to `default_module:default_program`
 - It's a Protelis-generated name for anonymous modules
 - In "Program", two entry appear: one is the round computation, the other one is the event sending information to neighbors
- You can use **P** to pause and restart the simulation

Speeding up the simulation in Alchemist

Improve the performance

- If you attached a profiler to Alchemist, you'll see that most of the time is spent in rendering nodes
- By default, the simulator renders the entire scene (within the viewport) after every event
 - This can get very expensive very quickly
- You can tune the refresh rate of the UI by pressing:
 - **Right arrow**: speed up (less calls to the graphics)
 - **Left arrow**: speed down (more calls to the graphics)
- If you speed up too much, the UI will become sluggish (e.g. if you update it only twice per second)
- **R** can be used to try to make the time flow linearly
 - Useful in case of sparse, unevenly distributed events
 - Left and right arrows keep working
- Lots of improvements are down the line, but not there yet :(

Loading a Protelis file I

- Writing Protelis specification within a YAML file is not what we want
 - No portability
 - Easy to make mistakes
 - Horrible kludges to make YAML happy with pieces of Protelis inside
 - Maybe you want to run your code elsewhere, not just inside a simulator
- Let's load a Protelis program from its own file!

tutorial/zero.pt

```
// Make sure that the module is in the classpath!
module tutorial:zero
0 // Program script
```

Loading a Protelis file II

03-load-module.yml

```
incarnation: protelis
network-model:
  type: EuclideanDistance
  parameters: [10]
displacements:
- in:
  type: Circle
  parameters: [500, 0, 0, 50]
programs:
-
- time-distribution: 1
  program: tutorial:zero # Just use the name of the module! The loading
    ↳ happens transparently from the classpath (which, of course, means
    ↳ that your Protelis source folder should be included in your Java
    ↳ classpath)
- program: send
```

- Run with: `./gradlew -Pfile=03-load-module`

Outline

1 Introduction

- Aggregate programming: idea and motivation
- Protelis as practical AP language and workbench setup

2 Core language constructs

- The boring stuff: functions, operators, and the like
- Stateful computation with `rep`
- Sharing information with `nbr`
- Domain restriction with `if`

3 Libraries

- Frequently used patterns: towards reusable building blocks
- Filling the abstraction gap: Protelis-lang

4 More about the toolchain

- Making nodes move
- Configuring batches and exporting data
- Using real world maps
- Getting from simulations to devices

Factorial in Protelis I

tutorial/factorial.pt

```
module tutorial:factorial
/*
 * The language is functional, every expression has a return value. In case of multiple statements in a
 * block, the value of the last expression is returned.
 *
 * Comments are C-like, both single and multiline supported.
 *
 * The following defines a new function. If the optional "public" keyword is present, the function will be
 * accessible from outside the module
 */
public def factorial(n) { // Dynamic typing
    if (n <= 1) {
        1 // No return keyword, no ";" at the end of the last line
    } else { // else is mandatory
        n * factorial(n - 1) // infix operators, recursion
    }
}
// There is no main function, just write the program at the end (Python-like)
let num = 5; // mandatory ";" for multiline instructions
factorial(5) // Function call
```

Factorial in Protelis II

04-factorial.yml

```
incarnation: protelis
network-model:
  type: EuclideanDistance
  parameters: [10]
displacements:
- in:
    type: Circle
    parameters: [500, 0, 0, 50]
programs:
-
- time-distribution: 1
  program: tutorial:factorial
- program: send
```

- Run with: `./gradlew -Pfile=04-factorial`

Factorial in Protelis III

Expected result

- All network nodes independently compute the result
- Those which computed get a purple circle around
- The value of the computation gets written aside them
 - Again, no magic. I configured an effect to do so.
 - Mastering the effect system of Alchemist is not trivial...
 - and maybe not worth doing it, as it will get updated soon
 - ...but may provide nice insights on the system
 - and nice screenshots on your papers

Higher order, lambdas, tuples, Java interoperability I

tutorial/randomfactorial.pt

```
module tutorial:randomfactorial
import java.lang.Math.random // Import java static methods
def dice() {
    let rand = random(); // call methods as if they were local functions
    (6 * rand).intValue() // Java-style method invocation on target objects
}
def factorial(n) {
    if (n <= 1) { 1 } else { n * factorial(n - 1) }
}
def callWithDice(fun) { // Higher-order function
    let num = dice();
    fun.apply(num)
}
[ // Use square brackets to build a tuple
    callWithDice(factorial), // higher-order call
    callWithDice((n) -> {n ^ 2}) // as above, but with an anonymous function (square)
] * [ 1, 2 ] // Tuple can be used with operators (they must have the same length though)
/* Elements of a tuple can be accessed using get(), several methods allow to create new tuples using
   ↪ existing ones. Refer to this javadoc: http://protelis-doc.surge.sh/ */
```

Higher order, lambdas, tuples, Java interoperability II

05-higherorder.yml

```
incarnation: protelis
network-model:
  type: EuclideanDistance
  parameters: [10]
displacements:
  - in:
      type: Circle
      parameters: [500, 0, 0, 50]
programs:
  -
    - time-distribution: 1
      program: tutorial:randomfactorial
    - program: send
```

- Run with: `./gradlew -Pfile=05-higherorder`

Accessing the device, sensing, actuating I

tutorial/deviceaccess.pt

```
module tutorial:deviceaccess
/*
 * Two important entities:
 * - self: provides access to the device implementation. Different devices have different capabilities, so
 *   some Protelis code may not be able to run on some platforms (e.g. code requiring access to the
 *   current coordinates may not work in a device implementation that does not expose such information).
 *   See ExecutionContext and its sub-interfaces at
 *   http://protelis-doc.surge.sh/org/protelis/vm/ExecutionContext.html
 * - env: provides access by name to sensors/actuators. They are treated uniformly as variables, similarly
 *   as accessing a map
 */
let injected = if (env.has("value")) { env.get("value") } else { "No value" }; // Read a value, if present
env.put("a generated value",
    (100 * self.nextRandomDouble().intValue()); // This is the best way to ask the platform for a random
self.getDeviceUID().getId() // The unique identifier of this device
```

Accessing the device, sensing, actuating II

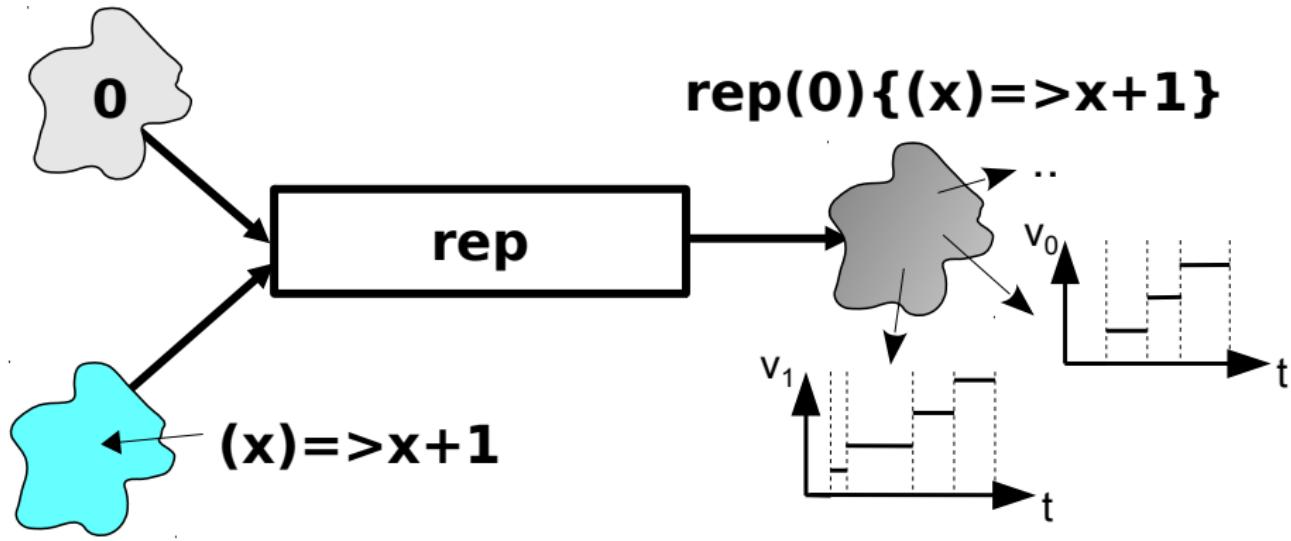
06-deviceaccess.yml

```
incarnation: protelis
network-model:
  type: EuclideanDistance
  parameters: [10]
displacements:
- in:
    type: Circle
    parameters: [500, 0, 0, 50]
programs:
-
- time-distribution: 1
  program: tutorial:deviceaccess
- program: send
contents: #in this section we can pre-configure the value of sensors
- in: # Optionally, we can limit the area of space where to inject such value
  type: Rectangle
  parameters: [0, 0, 100, 100]
molecule: value # name of the sensor
concentration: "somevalue" # initial value (any type)
```

- Run with: ./gradlew -Pfile=06-deviceaccess

Evolving fields with time I

- We only computed in a stateless fashion
 - No, using env as a mean to retain state is not a good idea
- In field calculus, **evolving fields** are built with the rep operator



Evolving fields with time II

Syntax and semantics in Protelis

```
rep (value <- initial) { body }
```

- When `rep` is encountered for the first time, `value` is set to `initial`
- Then `body` is evaluated, and the result stored in `value`
- So, in the next round, `value` will have the result of the evaluation of `body` in the previous round

A reset counter I

primitives/counter.pt

```
module primitives:counter
public def cyclicCounter(min, max) {
// Let's start from min - 1, so that at the first round we get min.
    rep(count <- min - 1) {
        // mux is a functional multiplexer: evaluates both branches and returns one
        // of them
        mux(count >= max) { // Here == would suffice
            0
        } else {
            count + 1
        }
    }
}
cyclicCounter(0, 10)
```

A reset counter II

07-counter.yml

```
incarnation: protelis
network-model:
  type: EuclideanDistance
  parameters: [10]
displacements:
  - in:
    type: Circle
    parameters: [500, 0, 0, 50]
    programs:
      -
        - time-distribution: 1
          program: primitives:counter
        - program: send
contents:
```

- Run with: `./gradlew -Pfile=07-counter`

Using different time distributions

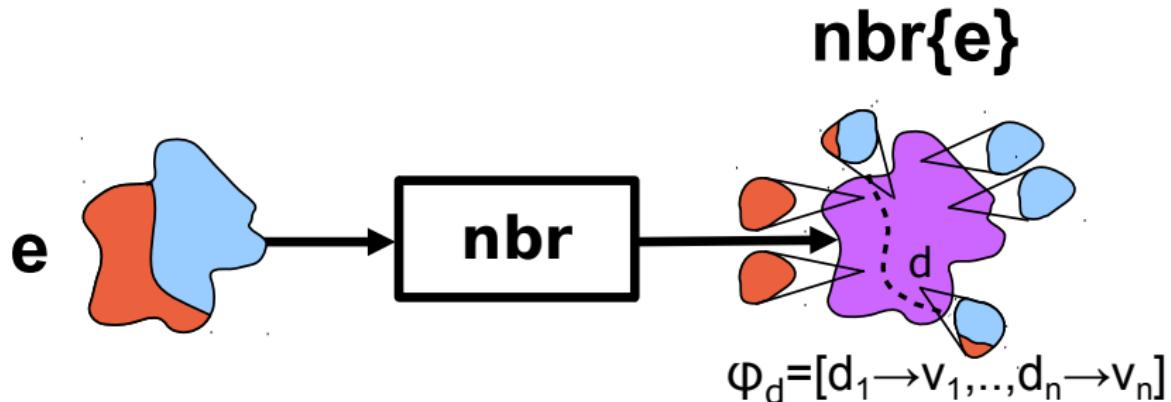
08-async.yml

```
incarnation: protelis
network-model:
  type: EuclideanDistance
  parameters: [10]
displacements:
  - in:
    type: Circle
    parameters: [500, 0, 0, 50]
  programs:
    -
      - time-distribution: # Check out in the Alchemist Javadocs all the classes implementing
        ↪ "TimeDistribution" if you want a complete overview of what's available
        type: ExponentialTime # Markovian-distributed events with rate = 1
        parameters: [1]
        program: primitives:counter
      - program: send
contents:
```

- Run with: `./gradlew -Pfile=08-async`
- Devices run with a random (and memoryless) time distribution (variance equal to the mean): very chaotic behavior.

Distribution: view of the field in the local neighborhood I

- We did not yet use the communication abilities of our devices
 - It's time to do it!
- In field calculus, **fields of fields** are built with the `nbr` operator
 - It builds a field of views of a field in the neighborhood



Distribution: view of the field in the local neighborhood II

Syntax and semantics in Protelis

`nbr(body)`

- Evaluates body
- Returns a field containing the mapping between the neighbouring *aligned* devices (included self) and the value of body at their location
 - A device is d_0 aligned to another device d_1 if, recently enough, d_0 executed the same `nbr` instruction d_1 is executing, sharing the result with it.
- Shares the local value of body

Reduction and manipulation of fields of fields I

Reducing fields

- The information on a field is usually reduced to a single value (a field, not a field of fields)
- Such operation is performed by the hood built-in function

```
hood(reduction, defval, fieldExpr)
```

```
hood PlusSelf(reduction, defval, fieldExpr)
```

Where reduction is a function of two parameters ($T, T \rightarrow T$), defval is a value of type T , and fieldExpr a field of fields of T

- fieldExpr is evaluated
- if PlusSelf is not specified, the local value is discarded
- reduction is applied to every value
- If there is a result it is returned, defval is returned otherwise

Reduction and manipulation of fields of fields II

Built-in hoods

Protelis has a number of hood operations ready to use (every one with its PlusSelf variant)

- minHood, minimum of a field, default Infinity
- maxHood, maximum of a field, default -Infinity
- anyHood, “or” on a field of booleans, default false
- allHood, “and” on a field of booleans, default true
- sumHood, sum of the values of the field, default 0
- meanHood, mean of the values of the field, default NaN
- unionHood, a tuple with all the values in the field, default []

Reduction and manipulation of fields of fields III

Manipulation of fields of fields

- Fields and locals (using the global view: fields of fields and fields) can be used in any operation
- e.g. `nbr(v) + 1` will produce as result a field containing for every neighbor the value of `v + 1`
- e.g. `nbr(v) + nbr(a)` will combine the values neighbor-per-neighbor
- it also works for method call: `nbr(anObject).hashCode()` will return a field of hash codes

Reduction and manipulation of fields of fields IV

Manipulation and alignment

- Manipulation of fields of fields is possible because the two instructions that can “break the alignment” (`rep` and `if`—to be seen soon) cannot return fields of fields
- Fields of fields are guaranteed to always have the same devices at every step of the execution!

Pace-making counter I

primitives/pacemaker.pt

```
module primitives:pacemaker
import primitives:counter // import the public function from another Protelis module
def max(a, b) { // a max working on anything that is comparable
    if (a > b) { a } else { b }
}
def counter() { rep (i <- -1) { i + 1 } } // plain counter
public def pacemaker(min, max) {
    let myCount = [counter(), cyclicCounter(min, max)]; // tuple with the count of my rounds and the value of
    ↪ the cyclic counter
    rep (count <- myCount) { // pick and store...
        maxHood PlusSelf(nbr(max(myCount, count))) // ...the maximum in my neighborhood between the maximum of
        ↪ my count and the maximum i got until now
    }.get(1) // we are interested in the second argument of the tuple, the first an implementation detail
}
pacemaker(0, 10)
```

Pace-making counter II

09-pacemaker.yml

```
incarnation: protelis
network-model:
  type: EuclideanDistance
  parameters: [10]
displacements:
- in:
  type: Circle
  parameters: [500, 0, 0, 50]
programs:
-
- time-distribution:
  type: ExponentialTime
  parameters: [1]
  type: Event # We define the event manually...
  actions: # ...because we want to specify a parameter of the Protelis interpreter.
  - type: RunProtelisProgram
    parameters: ["primitives:pacemaker", 5] # We discard messages after 5 simulated seconds
- program: send
```

- Run with: `./gradlew -Pfile=09-pacemaker`
- The fastest device generates a “wave” that pulls up the slow ones!

Gradient: computing distances I

primitives/gradient.pt

```
module primitives:gradient

/*
 * Compute the distance to the closest source. The input is a field of booleans, this function computes the
 * distance from the closest device where the field yields true.
 */
public def gradient(source) {
    rep (d <- Infinity) { // Start from an infinite distance
        mux (source) { // Functional multiplexer (evaluates both branches)
            0 // If I am the source, my distance is zero
        } else {
            minHood(nbr(d) + self.nbrRange()) // The distance is computed as the minimum of the sum of the
                // distance between me and my neighbors, and my neighbors and the source, for each neighbor.
        }
    }
    // Let's compute the distance from the device "zero"
    gradient(self.getDeviceUID().getId() == 0)
```

Gradient: computing distances II

10-gradient.yml

```
incarnation: protelis
network-model:
  type: EuclideanDistance
  parameters: [10]
displacements:
- in:
  type: Circle
  parameters: [500, 0, 0, 50]
  programs:
  -
    - time-distribution: 1
      program: primitives:gradient
    - program: send
```

- Run with: `./gradlew -Pfile=10-gradient`
- The fastest device generates a “wave” that pulls up the slow ones!

Gradient: computing distances I

primitives/gradient.pt

```
module primitives:gradient

/*
 * Compute the distance to the closest source. The input is a field of booleans, this function computes the
 * distance from the closest device where the field yields true.
 */
public def gradient(source) {
    rep (d <- Infinity) { // Start from an infinite distance
        mux (source) { // Functional multiplexer (evaluates both branches)
            0 // If I am the source, my distance is zero
        } else {
            minHood(nbr(d) + self.nbrRange()) // The distance is computed as the minimum of the sum of the
                // distance between me and my neighbors, and my neighbors and the source, for each neighbor.
        }
    }
    // Let's compute the distance from the device "zero"
    gradient(self.getDeviceUID().getId() == 0)
```

Gradient: computing distances II

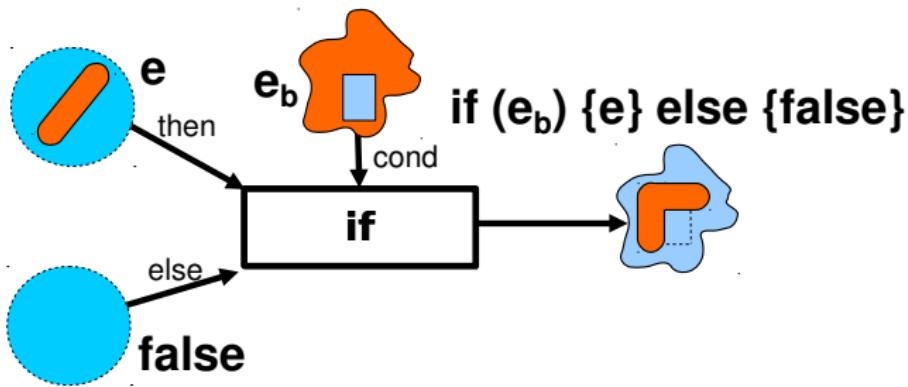
10-gradient.yml

```
incarnation: protelis
network-model:
  type: EuclideanDistance
  parameters: [10]
displacements:
- in:
  type: Circle
  parameters: [500, 0, 0, 50]
  programs:
  -
    - time-distribution: 1
      program: primitives:gradient
    - program: send
```

- Run with: `./gradlew -Pfile=10-gradient`
- Move the source node around and see what happens!

Domain restriction with if |

- In aggregate programming, using if has different consequences than the languages you are used to
 - It's not a "flow control" instruction
 - It's a **domain restriction** instruction
- Devices get separated in two, non communicating, domains
 - They lose *alignment*
 - Reasoning at the device level can be very misleading!



Domain restriction with if ||

Syntax and semantics in Protelis

```
if(condition) { then } else { otherwise }
```

- Evaluates condition
- Executes either then or otherwise, depending on the value of condition
- Unexecuted branches lose their state
 - reps get reset to default
- If one of the branches contains:
 - nbr
 - rep
 - a call to a function using the two above
- then the choice between if and mux is paramount!

Screw up the gradient I

primitives/brokengradient.pt

```
module primitives:brokengradient
def gradient(source) {
    rep (d <- Infinity) {
        if (source) { // Changed mux with if
            0
        } else {
            // The source never enters this domain!
            // Frome the single device perspective, it never executes this branch, and so it never shares its
            // knowledge of being a source!
            minHood(nbr(d) + self.nbrRange())
        }
    }
    gradient(self.getDeviceUID().getId() == 0)
```

Screw up the gradient II

11-brokengradient.yml

```
incarnation: protelis
network-model:
  type: EuclideanDistance
  parameters: [10]
displacements:
- in:
  type: Circle
  parameters: [500, 0, 0, 50]
  programs:
  -
    - time-distribution: 1
      program: primitives:brokengradient
    - program: send
```

- Run with: `./gradlew -Pfile=11-brokengradient`
- Our toy got bricked!

Mixing rep, if and mux !

primitives/twocounters.pt

```
module primitives:twocounters
def counter() { rep(c <- 0) { c + 1 } }
[
  if (counter() % 2 == 0) { counter() } else { counter() },
  mux (counter() % 2 == 0) { counter() } else { counter() }
]
```

12-twocounters.yml

```
incarnation: protelis
network-model:
  type: EuclideanDistance
  parameters: [10]
displacements:
  - in:
      type: Circle
      parameters: [500, 0, 0, 50]
      programs:
        -
          - time-distribution: 1
            program: primitives:twocounters
          - program: send
```

Mixing rep, if and mux II

- Run with: `./gradlew -Pfile=12-twocounters`
- if continuously resets the rep status!

Gradient in a sub-domain I

primitives/gradobs.pt

```
module primitives:gradobs
import primitives:gradient
if (env.has("obstacle")) {
    Infinity
} else {
    gradient(self.getDeviceUID().getId() == 0)
}
```

Gradient in a sub-domain II

13-gradobs.yml

```
incarnation: protelis
network-model:
  type: EuclideanDistance
  parameters: [10]
displacements:
- in:
  type: Circle
  parameters: [500, 0, 0, 50]
programs:
-
- time-distribution: 1
  program: primitives:gradobs
- program: send
contents:
- in:
  type: Rectangle
  parameters: [-15, -50, 30, 75]
molecule: obstacle
concentration: true
```

- Run with: `./gradlew -Pfile=13-gradobs`
- Modify the program to use mux, and see the difference!

Outline

1 Introduction

- Aggregate programming: idea and motivation
- Protelis as practical AP language and workbench setup

2 Core language constructs

- The boring stuff: functions, operators, and the like
- Stateful computation with `rep`
- Sharing information with `nbr`
- Domain restriction with `if`

3 Libraries

- Frequently used patterns: towards reusable building blocks
- Filling the abstraction gap: Protelis-lang

4 More about the toolchain

- Making nodes move
- Configuring batches and exporting data
- Using real world maps
- Getting from simulations to devices

Raise the abstraction level

- Programming in terms of `rep`, `nbr` and `if` is very, very low level
 - And also, very much error-prone
 - What would happen if `minHood PlusSelf` was used in the gradient?
 - Try it in the simulator, modifying the 10-gradient example
 - Move nodes around: it's not self-stabilizing anymore...
- Several coordination patterns are common to many algorithms
 - Spread information
 - Accumulate information
 - Evolve state with time
 - Partition the system
- We want to program with them directly
- There are several ways of implementing such primitives
 - With different properties, both functional and non-functional

The G / C / T / S building blocks I

G: spreading information

`G(source, initial, metric, accumulate)`

- Spreads initial from source
- initial can be of any type (T)
- source is a boolean
- metric is a function ()->Field<Number> that computes the gradient variation (the same role played by nbrRange in our gradient)
- values are combined with accumulate, a function (T)->T

Example: rewrite the gradient with G

```
def distanceTo(source) {  
    let metric = () -> { self.nbrRange() };  
    G(source, 0, metric, v -> { v + metric.apply() })  
}
```

The G / C / T / S building blocks II

C: accumulate information

`C(potential, reduce, local, null)`

- Accumulates local values along a potential
- potential is a Comparable type (typically a number)
- reduce is a $(T, T) \rightarrow T$
- local is the local value (T)
- null is the default value, in case there is nothing to accumulate

Example: Count the devices in the system, accumulating the information in source

```
def countDevices(source) {  
    C(distanceTo(source), sum, 1, 0)  
}
```

The G / C / T / S building blocks III

T: evolve state

`T(initial, zero, decay)`

- A timer from `initial` to `zero` with a decay rate.

S: partition the network

`S(grain, metric)`

- Elects a leader every grain, using the provided metric.
- Returns a boolean (true if the node is a leader)
- `grain` is a number with the partition size
- `metric` is a `()->Field<Number>` that computes the distances
- A timer from `initial` to `zero` with a decay rate.

Multiple gradients I

advanced/voronoi.pt

```
module advanced:voronoi
import protelis:coord:sparsechoice
import protelis:coord:spreading
let metric = () -> { self.nbrRange() };
G(S(25, metric), 0, metric, v -> { v + metric.apply() })
```

- It's one line of "true" code

Multiple gradients II

14-voronoi.yml

```
incarnation: protelis
network-model:
  type: EuclideanDistance
  parameters: [10]
displacements:
- in:
  type: Circle
  parameters: [500, 0, 0, 50]
programs:
-
- time-distribution: 1
  program: advanced:voronoi
- program: send
```

- Run with: `./gradlew -Pfile=14-voronoi`
- Try to move the nodes around, and see how the system adapts!

Limitations of the building block

- Programming with building blocks is way easier than fiddling with `reps` and `nbrs` by hand
- ...but it is not yet where we'd like to stand!
- Many “useful” behaviors can't be captured
 - as they are not self-stabilizing
 - Yet, they can be of use in many situations
 - Think of gossip, for instance
- It's very hard to get a grasp of the global situation
 - If there is an information spread from multiple sources, only information from the closest would make it
 - This is ok in many situations
 - But not enough in others
- Many algorithms recur very often
 - They should stay in a library

Use protelis-lang and focus on the application code

protelis-lang is a library¹ intended to:

- Provide a collection of the most widely used distributed algorithms
 - Contains over a hundred of them, from literature and other tools
 - If you develop something new, contribution is very welcome :)
- Shift the focus of the designer entirely on application code
- Provide means to overcome the limitations of the building blocks
 - Through a number of meta-algorithms that allow parallel aligned executions
 - Mainly by means of the novel `multiInstance` algorithm

protelis-lang is part of the Protelis distribution and ready to use

¹First introduced this year at a SASO Workshop: the 2nd edition of eCAS — did you follow the presentation? :)

Partition the network and spread multiple gradients I

advanced/voronoi.pt

```
module advanced:voronoi
import protelis:coord:sparsechoice
import protelis:coord:spreading
let metric = () -> { self.nbrRange() };
G(S(25, metric), 0, metric, v -> { v + metric.apply() })
```

- It's one line of "true" code

Partition the network and spread multiple gradients II

advanced/summarize.pt

```
module advanced:summarize
import protelis:lang:utils
import protelis:coord:spreading
import protelis:coord:sparsechoice
import protelis:coord:accumulation
import protelis:state:time
// Let's simulate a sensor measuring the temperature (in Celsius, range 20 to 30) every minute
env.put("temp", rep (v <- 25) { if(cyclicTimer(60)) { 20 + self.nextDouble() * 10 } else { v } });
// Let's compute the mean temperature of the whole network and share th infomation to all devices
let sink = S(Infinity, nbrRange);
summarize(sink, sum, env.get("temp"), 0) / broadcast(sink, countDevices(distanceTo(sink)))
```

Partition the network and spread multiple gradients III

15-summarize.yml

```
incarnation: protelis
network-model:
  type: EuclideanDistance
  parameters: [10]
displacements:
- in:
  type: Circle
  parameters: [500, 0, 0, 50]
  programs:
  -
    - time-distribution: 1
      program: advanced:summarize
    - program: send
```

- Run with: `./gradlew -Pfile=15-summarize`

Outline

1 Introduction

- Aggregate programming: idea and motivation
- Protelis as practical AP language and workbench setup

2 Core language constructs

- The boring stuff: functions, operators, and the like
- Stateful computation with `rep`
- Sharing information with `nbr`
- Domain restriction with `if`

3 Libraries

- Frequently used patterns: towards reusable building blocks
- Filling the abstraction gap: Protelis-lang

4 More about the toolchain

- Making nodes move
- Configuring batches and exporting data
- Using real world maps
- Getting from simulations to devices

Brownian movement I

16-brownian.yml

```
incarnation: protelis
network-model: { type: EuclideanDistance, parameters: [10] } # Compact syntax
program-pools:
  - compute-gradient: &gradient #This is a YAML anchor: assigns a name to this YAML object, and can be
    ↪ reused later in the file. It's not an Alchemist thing, it's pure YAML specification
    - time-distribution: 1
      program: primitives:gradient
    - program: send
  - move: &move
    - time-distribution: { type: ExponentialTime, parameters: [1] }
      type: Event
      actions:
        - { type: BrownianMove, parameters: [1] }
displacements:
  - in: { type: Circle, parameters: [500, 0, 0, 50] }
programs:
  # Now maybe the strange double list in "programs" makes sense :)
  - *gradient # Reference to the anchor previously defined.
  - *move
```

- Run with: `./gradlew -Pfile=16-brownian`

Controlling the movement direction I

advanced/converge.pt

```
module advanced:converge
import protelis:coord:sparsechoice
import protelis:coord:spreading
import protelis:state:time
let leader = S(25, nbrRange);
let destination = broadcast(leader, self.getCoordinates());
destination = if (isSignalStable(destination, 10)) { destination } else { self.getCoordinates() };
env.put("destination", destination);
self.getCoordinates() == destination
```

Controlling the movement direction II

17-converge.yml

```
incarnation: protelis
network-model: { type: EuclideanDistance, parameters: [10] }
program-pools:
  - compute-gradient: &gradient
    - time-distribution: 1
      program: advanced:converge
    - program: send
  - move: &move
    - time-distribution: { type: ExponentialTime, parameters: [1] }
      type: Event
      actions:
        - { type: MoveToTarget, parameters: [destination, 1] }
displacements:
  - in: { type: Circle, parameters: [500, 0, 0, 50] }
programs:
  - *gradient
  - *move
```

- Run with: `./gradlew -Pfile=17-converge`

Exporting data I

18-export.yml

```
incarnation: protelis
network-model: { type: EuclideanDistance, parameters: [10] }
program-pools:
  - compute-gradient: &gradient
    - { time-distribution: 1, program: "advanced:converge" }
    - program: send
  - move: &move
    - time-distribution: { type: ExponentialTime, parameters: [1] }
      type: Event
    actions: [ { type: MoveToTarget, parameters: [destination, 1] } ]
displacements: [ { in: { type: Circle, parameters: [500, 0, 0, 50] }, programs: [*gradient, *move] } ]
export:
  - time # Exports the current time
  - number-of-nodes # Exports the number of nodes in the system
  - molecule: advanced:converge # The name of the sensor / actuator whose value will be exported
    # Optionally, a "property: program" section can be specified, where "program" is a valid chunk of
    # protelis code. The sensor value will be processed using such code before being exported.
  value-filter: onlyfinite # You may want to keep poisonous values (Infinities and NaN) from being
    # exported or passed to the aggregator. Available filters are "nofilter" (default), "onlyfinite"
    # (discards both NaN and Infinities), "filternan", "filterinfinity".
aggregators: [sum, mean, kurtosis] # The aggregator takes all the data from nodes and reduces it to a
  # single value. If omitted, one value per node will be exported (in this case, 500 columns...). You
  # can load any UnivariateStatistics from the Apache Commons Math library, just by their name.
```

Exporting data II

- Run with: `./gradlew -Pfile=18-export`
- Look at the file `exported-data.txt`

Using variables and controlling randomness I

19-variables.yml

```

incarnation: protelis
variables:
  seed: &seed # You can give the anchor any name, assigning the name of the variable is convenient, though
    {min: 0, max: 9, step: 1, default: 0} # This variable ranges in [0, 9], steps of 1, defaulting to 0
  connection-radius: &connection-radius
    formula: 10 # An instance of a Javascript interpreter is created to evaluate the expression.
  round-frequency: &round-frequency
    formula: Math.max($seed, 1) # Other variables can be referred by prefixing them with $ (e.g. $seed would
      ↪ be substituted by the current value of seed).
  language: scala # The system defaults to Javascript, but a Scala interpreter can be instanced instead
  program-name: &program-name
    formula: "'advanced:converge'" # Use whatever type you want - just make sure you use them properly
seeds:
  scenario: *seed # This controls the initial displacement of the nodes
  simulation: *seed # This controls the generated event times
network-model: { type: EuclideanDistance, parameters: [*connection-radius] }
program-pools:
  - compute-gradient: &gradient
    - { time-distribution: 1, program: *program-name }
    - program: send
  - move: &move
    - time-distribution: { type: ExponentialTime, parameters: [*round-frequency] }
      type: Event
      actions: [ { type: MoveToTarget, parameters: [destination, 1] } ]
  displacements: [ { in: { type: Circle, parameters: [500, 0, 0, 50] }, programs: [*gradient, *move] } ]

```

Using variables and controlling randomness II

- Run with: `./gradlew -Pfile=19-variables`
- Run it again changing the default for the seed variable!
- Variables can be instanced with the type/parameters syntax
 - Several implementations available (geometrically scaling, arbitrary values...)
 - See <http://alchemist-doc.surge.sh/it/unibo/alchemist/loader/variables/package-frame.html>

Using variables and controlling randomness III

Batches

Independent variables (those with no formula entry) define the batch content

- if Alchemist is executed with the --batch, it will run in headless mode, spawning a simulation for each possible combination of values of the variables listed after the -var option
- e.g.: `java -jar alchemist.jar --batch --var seed range`
 - If range may yield 1, 10 or 100 and seed may yield 0 or 1
 - Six simulations would get generated
 - Don't abuse: with many variables and many variables things get big very quickly
- Simulations are executed in parallel
 - By default one simulation per core + 1
 - This behavior is configurable with a command line option
 - A grid executor is in our TODO list

Loading maps I

20-maps.yml

```
incarnation: protelis
environment:
  type: OSMEvironment
  parameters: ["/vcm.pbf"] # Loads from classpath, an absolute path to the file would work as well
pools:
  - pool: &move
    - time-distribution: 0.1
      type: Event
      actions:
        - type: ReproduceGPSTrace # Other strategies are available, such as interpolate the track using
          ↪ pedestrian roads
          parameters: ["/vcmuser.gpx", false, "AlignToTime", 1365922800, false, false]
displacements:
  - in:
    type: FromGPSTrace # displace nodes as they are in the GPS track
    parameters: [1497, "/vcmuser.gpx", false, "AlignToTime", 1365922800, false, false]
programs:
  - *move
```

- Run with: ./gradlew -Pfile=20-maps

Writing a backend I

Starting point: provide an ExecutionContext

Your back-end entry point is the implementation of the ExecutionContext interface

- Many complicated functions
- Fortunately, you can just extend from AbstractExecutionContext
 - Just a single `instance()` method left virtual, basically a call to the constructor.
 - You can see how's implemented in the simulator for a clue

Writing a backend II

Expose device capabilities

The `ExecutionContext` also holds device capabilities (it's like the `System` / `Runtime` singletons of Java).

- A number of interfaces for standard capabilities are provided
 - `TimeAwareDevice` — Implement this if your device is able to estimate network message delays and lags
 - `SpatiallyEmbeddedDevice` — Implement this if your device is able to estimate distances to its neighbors (with whatever metric)
 - `LocalizedDevice` — If your device has access to its position (e.g., if a GPS system is onboard)
- You are free to add any method, and access it via calls to `self`
- The thing usually implemented is something along the line of

```
MyExecutionContext extends AbstractExecutionContext implements
TimeAwareDevice, LocalizedDevice
```

Writing a backend III

Implement communication

Communication is dealt with by the NetworkManager

- `void shareState(Map<CodePath, Object> toSend)`
 - You are in charge of finding a way to send your neighbors the `toSend` payload...
- `Map<DeviceUID, Map<CodePath, Object>> getNeighborState()`
 - ...and of presenting the payloads sent to others to the current device.

Writing a backend IV

Run

Now that pieces are in place, just create a ProtelisVM and run whatever you like, e.g.:

```
final NetworkManager netmgr = new MyNetworkManager(...);
final ExecutionContext ctx = new MyExecutionContext(netmgr, ...);
final ProtelisProgram prog = ProtelisLoader.parse(myProgram)
final ProtelisVM vm = new ProtelisVM(prog, ctx);
while (alive) {
    vm.runCycle();
    Thread.sleep(someTime);
}
```

References I



Oren Ben-Kiki, Clark Evans, and Ingy döt Net.

Yaml ain't markup language (yaml™), version 1.2, 3rd edition, 2009-10-01.
<http://yaml.org/spec/1.2/spec.pdf>.



Jacob Beal, Danilo Pianini, and Mirko Viroli.

Aggregate programming for the internet of things.
IEEE Computer, 48(9):22–30, 2015.



Jacob Beal, Mirko Viroli, Danilo Pianini, and Ferruccio Damiani.

Self-adaptation to device distribution changes.

In *10th IEEE International Conference on Self-Adaptive and Self-Organizing Systems, SASO 2016, Augsburg, Germany, September 12-16, 2016*, pages 60–69, 2016.



Sara Montagna, Danilo Pianini, and Mirko Viroli.

Gradient-based self-organisation patterns of anticipative adaptation.

In *Sixth IEEE International Conference on Self-Adaptive and Self-Organizing Systems, SASO 2012, Lyon, France, September 10-14, 2012*, pages 169–174, 2012.

References II



Sara Montagna, Danilo Pianini, and Mirko Viroli.

A model for drosophila melanogaster development from a single cell to stripe pattern formation.

In *Proceedings of the ACM Symposium on Applied Computing, SAC 2012, Riva, Trento, Italy, March 26-30, 2012*, pages 1406–1412, 2012.



Sara Montagna, Mirko Viroli, Matteo Risoldi, Danilo Pianini, and Giovanna Di Marzo Serugendo.

Self-organising pervasive ecosystems: A crowd evacuation example.

In *Software Engineering for Resilient Systems - Third International Workshop, SERENE 2011, Geneva, Switzerland, September 29-30, 2011. Proceedings*, pages 115–129, 2011.



Danilo Pianini, Jacob Beal, and Mirko Viroli.

Improving gossip dynamics through overlapping replicates.

In *Coordination Models and Languages - 18th IFIP WG 6.1 International Conference, COORDINATION 2016, Held as Part of the 11th International Federated Conference on Distributed Computing Techniques, DisCoTec 2016, Heraklion, Crete, Greece, June 6-9, 2016, Proceedings*, pages 192–207, 2016.

References III



Danilo Pianini, Simon Dobson, and Mirko Viroli.

Self-stabilising target counting in wireless sensor networks using euler integration.

11th IEEE International Conference on Self-Adaptive and Self-Organizing Systems, SASO 2017, Tucson, AZ, USA, September 18-22, 2017.

To appear.



Danilo Pianini, Sara Montagna, and Mirko Viroli.

Chemical-oriented simulation of computational systems with ALCHEMIST.

J. Simulation, 7(3):202–215, 2013.



Danilo Pianini, Mirko Viroli, and Jacob Beal.

Protelis: practical aggregate programming.

In *Proceedings of the 30th Annual ACM Symposium on Applied Computing, Salamanca, Spain, April 13-17, 2015*, pages 1846–1853, 2015.



Graeme Stevenson, Juan Ye, Simon Dobson, Danilo Pianini, Sara Montagna, and Mirko Viroli.

Combining self-organisation, context-awareness and semantic reasoning: the case of resource discovery in opportunistic networks.

In *Proceedings of the 28th Annual ACM Symposium on Applied Computing, SAC '13, Coimbra, Portugal, March 18-22, 2013*, pages 1369–1376, 2013.

References IV



Mirko Viroli, Antonio Bucchiarone, Danilo Pianini, and Jacob Beal.
Combining self-organisation and autonomic computing in cass with aggregate-mape.
pages 186–191, 2016.